

# Restricted Boltzmann machine

Artificial neural networks

Christoffer Arvidsson

## Results

Parameters used to estimate the model distribution is shown in Table 1. Notably, the lower learning rate helped training the model for small number of hidden neurons.

Table 1: The parameters used in the model.

Parameter	Value
$k$	100
number of minibatches (trials)	3000
$n_{\text{visible}}$	3
$n_{\text{hidden}}$	{1,2,3,4,5,6,7,8}
learning rate $\eta$	0.001
minibatch size	20
generation realisations	2000
generation samples/realization	1000
averaging runs	3

The KL-divergence over number of hidden neurons is shown in Figure 1. We can see that three neurons seems to be the critical number of hidden neurons. These results do align with the theoretical upper bound, although the KL-divergence is slightly higher than this upper bound. This could be due to the CD-k algorithm converging to a suboptimal solution.

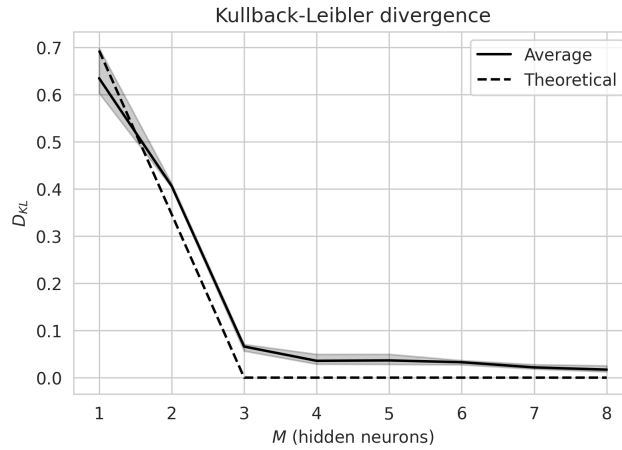


Figure 1: The KL-divergence over the number of hidden neurons in the Boltzmann machine for the XOR problem. The shaded region is the confidence interval. The dashed line is the theoretical upper bound.

# Notebook

## RBM implementation

```
1 import numpy as np
2
3 class RBM:
4     def __init__(self, n_visible, n_hidden):
5         self.weights = np.random.normal(0, 1, size=(n_hidden, n_visible))
6         self.t_vis = np.zeros((n_visible,), dtype=float)
7         self.t_hid = np.zeros((n_hidden,), dtype=float)
8
9     def generate(self, x, num_samples):
10        v_state = x.copy()
11        h_state = np.zeros_like(self.t_hid, dtype=int)
12
13        # update hidden neurons
14        b_h0 = np.dot(v_state, self.weights.T) - self.t_hid
15        p_b = (1+np.exp(-2*b_h0))**-1
16        r = np.random.uniform(size=h_state.shape)
17        h_state[r<p_b] = 1
18        h_state[r>=p_b] = -1
19
20        patterns = np.zeros((num_samples, x.shape[0]), dtype=int)
21
22        for i in range(num_samples):
23            # update visible neurons
24            b_v = np.dot(h_state, self.weights) - self.t_vis
25            p_b = (1+np.exp(-2*b_v))**-1
26            r = np.random.uniform(size=v_state.shape)
27            v_state[r<p_b] = 1
28            v_state[r>=p_b] = -1
29
30            # update hidden neurons
31            b_h = np.dot(v_state, self.weights.T) - self.t_hid
32            p_b = (1+np.exp(-2*b_h))**-1
33            r = np.random.uniform(size=h_state.shape)
34            h_state[r<p_b] = 1
35            h_state[r>=p_b] = -1
36
37            patterns[i] = v_state
38
39        return patterns
40
41    def run_cd_k(self, batch, k=100, learning_rate=0.1):
42        dW = np.zeros_like(self.weights, dtype=float)
43        dT_vis = np.zeros_like(self.t_vis, dtype=float)
44        dT_hid = np.zeros_like(self.t_hid, dtype=float)
45        h_state = np.zeros_like(self.t_hid, dtype=int)
46
47        for x in batch:
48            v_state = x.copy()
49
50            # update hidden neurons
51            b_h = np.dot(v_state, self.weights.T) - self.t_hid
52            b_h0 = b_h.copy()
53            p_b = (1+np.exp(-2*b_h))**-1
54            r = np.random.uniform(size=h_state.shape)
55            h_state[r<p_b] = 1
56            h_state[r>=p_b] = -1
57
58            for t in range(k):
59                # update visible neurons
60                b_v = np.dot(h_state, self.weights) - self.t_vis
61                p_b = (1+np.exp(-2*b_v))**-1
62                r = np.random.uniform(size=v_state.shape)
63                v_state[r<p_b] = 1
64                v_state[r>=p_b] = -1
65
66                # update hidden neurons
67                b_h = np.dot(v_state, self.weights.T) - self.t_hid
68                p_b = (1+np.exp(-2*b_h))**-1
69                r = np.random.uniform(size=h_state.shape)
70                h_state[r<p_b] = 1
71                h_state[r>=p_b] = -1
72
73            dW += learning_rate*(np.outer(np.tanh(b_h0), x) - np.outer(np.tanh(b_h), v_state))
74            dT_vis -= learning_rate*(x - v_state)
75            dT_hid -= learning_rate*(np.tanh(b_h0) - np.tanh(b_h))
76
77        self.weights += dW
78        self.t_vis += dT_vis
```

```
79         self.t_hid += dT_hid
```

## Dataset

```
1 import numpy as np
2 import seaborn as sns
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 from model import RBM
6 from multiprocessing import Pool
7 import os
8
9 sns.set_style("whitegrid")
10 plt.rcParams['figure.dpi'] = 300
11 plt.rcParams['savefig.dpi'] = 300
12
13 dataset = np.array([
14     [-1,-1,-1],
15     [1,-1,1],
16     [-1,1,1],
17     [1,1,-1],
18 ])
```

## Simulation

Define the simulation function

```
1 def run_simulation(dataset, k, weight_updates, n_visible, n_hidden,
2                     learning_rate, batch_size,
3                     n_samples, n_realizations):
4
5     machine = RBM(n_visible, n_hidden)
6
7     # Training
8     mus = np.random.choice(dataset.shape[0], size=(weight_updates, batch_size), replace=True)
9     counts = np.zeros(dataset.shape[0])
10    for i in range(weight_updates):
11        mu = mus[i]
12        machine.run_cd_k(dataset[mu], k=k, learning_rate=learning_rate)
13
14    # Generation
15    random_patterns = np.random.choice([-1, 1], size=(n_realizations, n_visible),
16                                             replace=True)
17    patterns = np.zeros((n_realizations, n_samples, n_visible),
18                        dtype=int)
19    for r in range(n_realizations):
20        patterns[r] = machine.generate(random_patterns[r], n_samples)
21
22    patterns[patterns == -1] = 0
23    n_patterns = n_realizations*n_samples
24    patterns = patterns.reshape((n_patterns,-1))
25    hashes = patterns.dot(1 << np.arange(patterns.shape[-1]-1, -1, -1))
26    unique, counts = np.unique(hashes, return_counts=True)
27
28    distribution = np.zeros(2**n_visible)
29    distribution[unique] = counts
30
31    p_model = distribution/(n_patterns)
32
33    return machine, p_model
```

Parameters used.

```
1 k = 100 # monte carlo iterations
2 weight_updates = 3000
3 n_visible = 3 # N
4 n_hidden = np.arange(1,8+1)
5 learning_rate = 0.001
6 batch_size = 20
7 n_realizations = 2000
8 n_samples = 1000
9 num_processes = 12 # Number of threads you have access to on your cpu.
10 averaging_runs = 3
```

Let's do some multiprocessing to speed this up. This runs the simulation for each value of `n_visible`, averaging\_runs number of times.

```

1 def _pool_func(n_hidden):
2     print(f'pid: {os.getpid()}\tRunning for n_hidden: {n_hidden}\n')
3     return run_simulation(dataset, k, weight_updates, n_visible, n_hidden,
4                           learning_rate, batch_size,
5                           n_samples, n_realizations)
6
7 requests = []
8 p_models = np.zeros((averaging_runs, len(n_hidden), 2**n_visible), dtype=float)
9 with Pool(processes=num_processes) as pool:
10     for r in range(averaging_runs):
11         requests.append([pool.apply_async(_pool_func, (n,)) for n in n_hidden])
12
13     for r in range(averaging_runs):
14         results = [req.get() for req in requests[r]]
15         # Effectively transposes our list of tuples into a tuple of lists
16         _, p_model = map(list, zip(*results))
17         p_models[r] = np.array(p_model)
18
19 print('Done with simulations.')
```

Now plot the KL-divergence

```

1 xor = [0,3,5,6] # Represents indices for the xor patterns among random ones
2 p_data = np.zeros((2**n_visible))
3 p_data[xor] = 1/4
4
5 def kl_divergence_bound(n, m):
6     return np.log(2) * (n - np.floor(np.log2(m + 1)) - (m+1)/(2**(np.floor(np.log2(m+1)))))
7
8 samples = []
9 for r in range(averaging_runs):
10     for i, p_model in enumerate(p_models[r]):
11         kl_divergence = np.sum(p_data[xor] * np.log(p_data[xor] / p_model[xor]))
12         samples.append((r, n_hidden[i], kl_divergence))
13
14 nonzero_m = n_hidden[n_hidden < 2**(n_visible - 1) - 1]
15 theoretical = np.zeros(len(n_hidden))
16 theoretical[n_hidden < 2**(n_visible - 1) - 1] = kl_divergence_bound(n_visible, nonzero_m)
17
18 data = pd.DataFrame(samples, columns=['run', 'hidden_neurons', 'kl_divergence'])
19 g = sns.lineplot(x='hidden_neurons', y='kl_divergence', data=data, color='black', label='Average')
20 sns.lineplot(x=n_hidden, y=theoretical, ax=g, color='black', linestyle='--', label='Theoretical')
21 g.set_xlabel('$M$ (hidden neurons)')
22 g.set_ylabel(r'$D_{KL}$')
23 g.set_title('Kullback-Leibler divergence')
24 plt.show()
```

